

Fork()

create a child process

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t fork(void);
```

Process resource utilizations (`getrusage(2)`) and CPU time counters (`times(2)`) are reset to zero in the child.

The child process is created with a single thread—the one that called `fork()`. The entire virtual address space of the parent is replicated in the child, including the states of mutexes, condition variables, and other pthreads objects; the use of `pthread_atfork(3)` may be helpful for dealing with problems that this can cause.

After a `fork()` in a multithreaded program, the child can safely call only `async-signal-safe` functions (see `signal-safety(7)`) until such time as it calls `execve(2)`.

The child inherits copies of the parent's set of open file descriptors.

The child inherits copies of the parent's set of open message queue descriptors

The child inherits copies of the parent's set of open directory streams

On success, the PID of the child process is returned in the parent, and 0 is returned in the child.

On failure, -1 is returned in the parent, no child process is created.

ERRORS

the `RLIMIT_NPROC` soft resource limit

the kernel's system-wide limit on the number of processes and threads

the maximum number of PIDs

failed to allocate the necessary kernel structures because memory is tight.

An attempt was made to create a child process in a PID namespace whose "init" process has terminated.

`fork()` is not supported on this platform

System call was interrupted by a signal and will be restarted.

Exit()

cause normal process termination

```
#include <stdlib.h>
```

```
void exit(int status);
```

Note that a call to `execve(2)` removes registrations created using `atexit(3)` and `on_exit(3)`.

After `exit()`, the exit status must be transmitted to the parent process. There are three cases:

- If the parent has set `SA_NOCLDWAIT`, or has set the `SIGCHLD` handler to `SIG_IGN`, the status is discarded and the child dies immediately.
- If the parent was waiting on the child, it is notified of the exit status and the child dies immediately.

- Otherwise, the child becomes a "zombie" process: most of the process resources are recycled, but a slot containing minimal information about the child process (termination status, resource usage statistics) is retained in process table. This allows the parent to subsequently use `waitpid(2)` (or similar) to learn the termination status of the child; at that point the zombie process slot is released.

Wait()

`wait`, `waitpid`, `waitid` - wait for process to change state

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t wait(int *wstatus);
```

```
pid_t waitpid(pid_t pid, int *wstatus, int options);
```

```
int waitid(idtype_t idtype, id_t id, siginfo_t *infop, int options);
```

`waitpid`

The value of `pid` can be:

- < -1 meaning wait for any child process whose process group ID is equal to the absolute value of `pid`.
- 1 meaning wait for any child process.
- 0 meaning wait for any child process whose process group ID is equal to that of the calling process.
- > 0 meaning wait for the child whose process ID is equal to the value of `pid`.

The value of `options` is an OR of zero or more of the following constants:

WNOHANG return immediately if no child has exited.

WUNTRACED also return if a child has stopped (but not traced via `trace(2)`).

 Status for traced children which have stopped

 is provided even if this option is not specified.

WCONTINUED also return if a stopped child has been resumed by delivery

 of `SIGCONT`.

If `wstatus` is not `NULL`, `wait()` and `waitpid()` store status information in the `int` to which it points. This integer can be inspected with the following macros:

`WIFEXITED(wstatus)`

 returns true if the child terminated normally, that is, by calling `exit()` or by returning from `main()`.

`WEXITSTATUS(wstatus)`

 returns the exit status of the child.

`WIFSIGNALED(wstatus)`

 returns true if the child process was terminated by a signal.

`WTERMSIG(wstatus)`

 returns the number of the signal that caused the child process to terminate.

`WCOREDUMP(wstatus)`

 returns true if the child produced a core dump.

`WIFSTOPPED(wstatus)`

 returns true if the child process was stopped by delivery of a signal

`WSTOPSIG(wstatus)`

 returns the number of the signal which caused the child to stop.

`WIFCONTINUED(wstatus)`

 returns true if the child process was resumed by delivery of `SIGCONT`.

waitid()

The `waitid()` system call provides more precise control over which child state changes to wait for.

The `idtype` and `id` arguments select the child(ren) to wait for, as follows:

`idtype == P_PID`

Wait for the child whose process ID matches `id`.

`idtype == P_PGID`

Wait for any child whose process group ID matches `id`.

`idtype == P_ALL`

Wait for any child; `id` is ignored.

RETURN VALUE

`wait()`: on success, returns the process ID of the terminated child; on error, -1 is returned.

`waitpid()`: on success, returns the process ID of the child whose state has changed; if `WNOHANG` was specified and one or more child(ren) specified by `pid` exist, but have not yet changed state, then 0 is returned. On error, -1 is returned.

`waitid()`: returns 0 on success or if `WNOHANG` was specified and no child(ren) specified by `id` has yet changed state; on error, -1 is returned.

Execv()

execl, execlp, execl, execv, execvp, execvpe - execute a file

```
#include <unistd.h>
```

```
extern char **environ;
```

```
int execl(const char *path, const char *arg, ... /* (char *) NULL */);
```

```
int execlp(const char *file, const char *arg, ... /* (char *) NULL */);
```

```
int execl(const char *path, const char *arg, ... /*, (char *) NULL, char * const envp[] */);
```

```
int execv(const char *path, char *const argv[]);
```

```
int execvp(const char *file, char *const argv[]);
```

```
int execvpe(const char *file, char *const argv[], char *const envp[]);
```

The `const char *arg` and subsequent ellipses in the `execl()`, `execlp()`, and `execl()` functions can be thought of as `arg0`, `arg1`, ..., `argn`. Together they describe a list of one or more pointers to null-terminated strings that represent the argument list available to the executed program.

Special semantics for `execlp()` and `execvp()`

The `execlp()`, `execvp()`, and `execvpe()` functions duplicate the actions of the shell in searching for an executable file if the specified file name does not contain a slash (/) character. The file is sought in the colon-separated list of directory pathnames specified in the `PATH` environment variable. If this variable isn't defined, the path list defaults to a list that includes the directories returned by `conf-str(_CS_PATH)` and possibly also the current working directory.

If the specified filename includes a slash character, then `PATH` is ignored, and the file at the specified pathname is executed.

The `exec()` functions return only if an error has occurred. The return value is `-1`.