

Select()

SYNOPSIS

```
/* According to POSIX.1-2001, POSIX.1-2008 */

#include <sys/select.h>

/* According to earlier standards */

#include <sys/time.h>

#include <sys/types.h>

#include <unistd.h>

int select(int nfds, fd_set *readfds, fd_set *writefds,
           fd_set *exceptfds, struct timeval *timeout);

void FD_CLR(int fd, fd_set *set);

int  FD_ISSET(int fd, fd_set *set);

void FD_SET(int fd, fd_set *set);

void FD_ZERO(fd_set *set);


#include <sys/select.h>

int pselect(int nfds, fd_set *readfds, fd_set *writefds,
            fd_set *exceptfds, const struct timespec *timeout,
            const sigset_t *sigmask);
```

Feature Test Macro Requirements for glibc (see `feature_test_macros(7)`):

`pselect()`: `_POSIX_C_SOURCE` \geq 200112L

DESCRIPTION

`select()` and `pselect()` allow a program to monitor multiple file descriptors, waiting until one or more of the file descriptors become "ready" for some class of I/O operation. A file descriptor is considered ready if it is possible to perform a corresponding I/O operation.

`select()` can monitor only file descriptors numbers that are less than `FD_SET_SIZE`

The operation of `select()` and `pselect()` is identical, other than these three differences:

- (i) `select()` uses a timeout that is a `struct timeval` (with seconds and microseconds), while `pselect()` uses a `struct timespec` (with seconds and nanoseconds).
- (ii) `select()` may update the timeout argument to indicate how much time was left. `pselect()` does not change this argument.
- (iii) `select()` has no sigmask argument, and behaves as `pselect()` called with `NULL` sigmask.

RETURN VALUE

On success, `select()` and `pselect()` return the number of file descriptors contained in the three returned descriptor sets (that is, the total number of bits that are set in `readfds`, `writfds`, `exceptfds`) which may be zero if the timeout expires before anything interesting happens. On error, `-1` is returned, and `errno` is set to indicate the error; the file descriptor sets are unmodified, and timeout becomes undefined.

Poll()

SYNOPSIS

```
#include <poll.h>

int poll(struct pollfd *fds, nfds_t nfd, int timeout);

#define _GNU_SOURCE          /* See feature_test_macros(7) */

#include <signal.h>

#include <poll.h>

int ppoll(struct pollfd *fds, nfds_t nfd,
          const struct timespec *tmo_p, const sigset_t *sigmask);
```

DESCRIPTION

`poll()` performs a similar task to `select(2)`: it waits for one of a set of file descriptors to become ready to perform I/O.

The set of file descriptors to be monitored is specified in the `fds` argument,

```
struct pollfd {
    int    fd;          /* file descriptor */
    short  events;       /* requested events */
    short  revents;      /* returned events */
};
```

The caller should specify the number of items in the `fds` array in `nfd`.

ppoll()

The relationship between `poll()` and `ppoll()` is analogous to the relationship between `select(2)` and `pselect(2)`: like `pselect(2)`, `ppoll()` allows an application to safely wait until either a file descriptor becomes ready or until a signal is caught.

RETURN VALUE

On success, a positive number is returned; this is the number of structures which have nonzero `revents` fields (in other words, those descriptors with events or errors reported). A value of 0 indicates that the call timed out and no file descriptors were ready. On error, -1 is returned, and `errno` is set appropriately.

Epoll_wait()

SYNOPSIS

```
#include <sys/epoll.h>

int epoll_wait(int epfd, struct epoll_event *events, int maxevents, int timeout);

int epoll_pwait(int epfd, struct epoll_event *events, int maxevents, int timeout, const sigset_t
*sigmask);
```

The `epoll_wait()` system call waits for events on the `epoll(7)` instance referred to by the file descriptor `epfd`. The memory area pointed to by `events` will contain the events that will be available for the caller. Up to `maxevents` are returned by `epoll_wait()`. The `maxevents` argument must be greater than zero.

`epoll_pwait()`

The relationship between `epoll_wait()` and `epoll_pwait()` is analogous to the relationship between `select(2)` and `pselect(2)`: like `pselect(2)`, `epoll_pwait()` allows an application to safely wait until either a file descriptor becomes ready or until a signal is caught.

RETURN VALUE

When successful, `epoll_wait()` returns the number of file descriptors ready for the requested I/O, or zero if no file descriptor became ready during the requested timeout milliseconds. When an error occurs, `epoll_wait()` returns -1 and `errno` is set appropriately.